

New York City Bike Trip Analysis

Peyman Kor

12/12/2019

Contents

Initialization, Import Required Libraries	2
SQL for Query	2
DATA Overview	2
Question: What is the trip duration distribution of Citibike trips?	4
Question: What is the most popular Citibike trip?	10
Question: Were there new bike stations introduced or removed at any point in time? What makes you think it were or weren't?	13

Initialization, Import Required Libraries

```
library(bigrquery)
library(tidyverse)
```

Now, the next step of the analysis would be the writing the SQL query to retrieve the data from the Google Cloud Platform. The data set could be found from: [LINK TO DATA SET](#)

SQL for Query

Firs, in order to have the overview on the data-set, we import all the columns and but in order to avoid the massive data-set, we limit the rows to 1000.

```
# Project ID
project_id <- "xenon-muse-261616" # put your project ID here

# query
sql_string <- "SELECT *
FROM
  [bigrquery-public-data.new_york.citibike_trips]
LIMIT
  1000"

# Execute the query and store the result
NY_sample <- query_exec(sql_string, project = project_id,
                        useLegacySql = FALSE)
```

```
## Using an auto-discovered, cached token.
## To suppress this message, modify your code or options to clearly consent to the use of a cached token
## See gargle's "Non-interactive auth" vignette for more details:
## https://gargle.r-lib.org/articles/non-interactive-auth.html
## The bigrquery package is using a cached token for kor.peyman@gmail.com.

## 0 bytes processed
```

Now, have a look on the head of the data:

DATA Overview

First, let's have the first look on the head of the data to get to know the data.

```
head(NY_sample)
```

```
##   tripduration      starttime      stoptime start_station_id
## 1          766 2016-03-22 17:24:23 2016-03-22 17:37:09         3230
## 2         96235 2015-10-23 19:36:55 2015-10-24 22:20:51          493
## 3         36395 2016-03-15 21:36:55 2016-03-16 07:43:31          375
```

```
## 4      659210 2015-11-04 07:20:10 2015-11-11 22:27:00      519
## 5      139 2015-11-02 16:56:06 2015-11-02 16:58:25      311
## 6      178 2016-01-15 07:09:38 2016-01-15 07:12:36      311
##      start_station_name start_station_latitude start_station_longitude
## 1      Penn Station Valet      40.75128      -73.99692
## 2      W 45 St & 6 Ave      40.75680      -73.98291
## 3      Mercer St & Bleecker St      40.72679      -73.99695
## 4      Pershing Square North      40.75187      -73.97771
## 5      Norfolk St & Broome St      40.71723      -73.98802
## 6      Norfolk St & Broome St      40.71723      -73.98802
##      end_station_id      end_station_name end_station_latitude
## 1      3240 NYCBS Depot BAL - DYR      0.00000
## 2      3019      NYCBS Depot - DEL      40.71663
## 3      3019      NYCBS Depot - DEL      40.71663
## 4      3019      NYCBS Depot - DEL      40.71663
## 5      3019      NYCBS Depot - DEL      40.71663
## 6      3019      NYCBS Depot - DEL      40.71663
##      end_station_longitude bikeid      usertype birth_year gender
## 1      0.00000 22812 Subscriber      1975      male
## 2      -73.98193 18137 Subscriber      1957      male
## 3      -73.98193 19910 Subscriber      1979      male
## 4      -73.98193 20894 Subscriber      1965      male
## 5      -73.98193 15608 Subscriber      1993 female
## 6      -73.98193 19800 Subscriber      1988 female
```

Another important will be to look whether the data contain any NA values, so that we could be cautious about our analysis in the next steps. Now, the following code will provide the NA values in each columns:

```
colSums(is.na(NY_sample))
```

```
##      tripduration      starttime      stoptime
##      0      0      0
##      start_station_id      start_station_name start_station_latitude
##      0      0      0
## start_station_longitude      end_station_id      end_station_name
##      0      0      0
##      end_station_latitude      end_station_longitude      bikeid
##      0      0      0
##      usertype      birth_year      gender
##      0      139      0
```

Mostly GOOD!. However, we notice that the *birth_year* column contains NA values. Another point is the *Column Type* we want to look that the imported data has the right format as we believe based on our understanding of the data, let's have look:

```
sapply(NY_sample, class)
```

```
## $tripduration
## [1] "integer"
##
## $starttime
## [1] "POSIXct" "POSIXt"
```

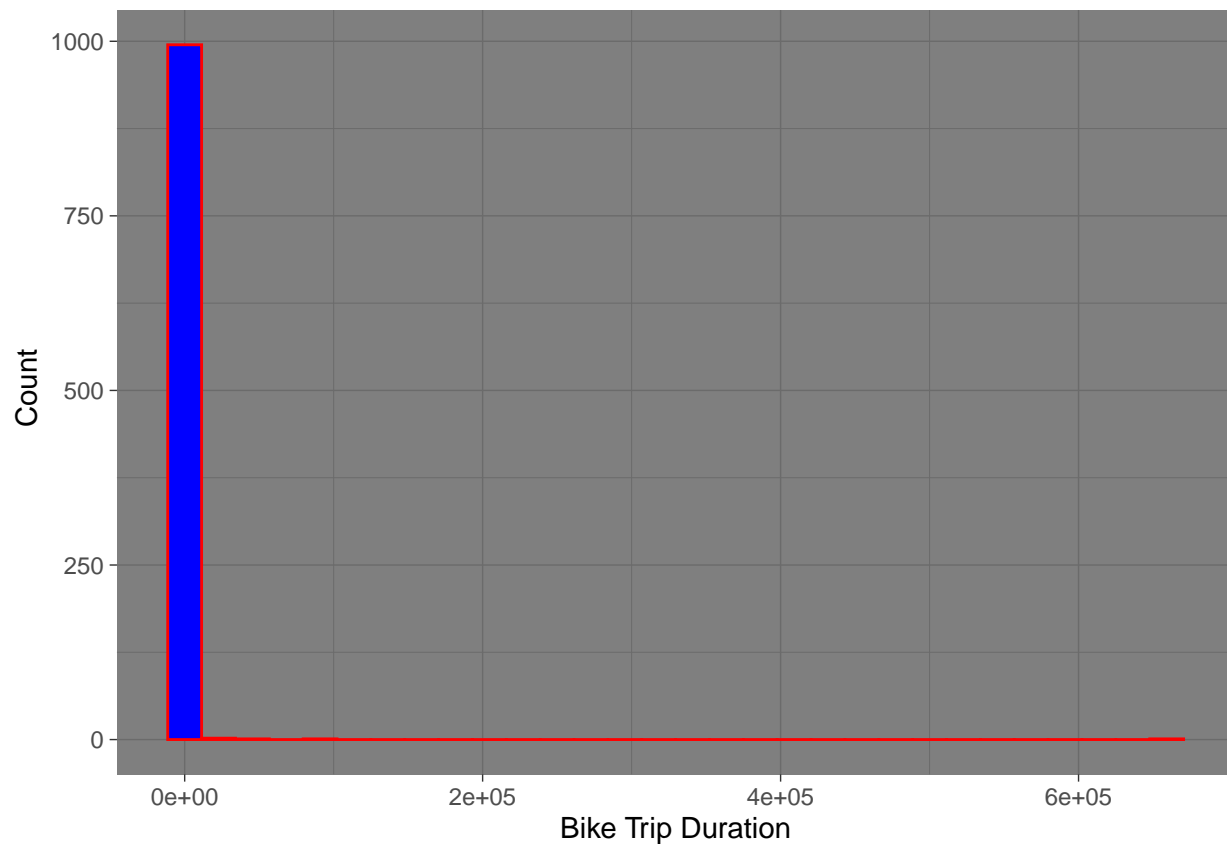
```
##
## $stoptime
## [1] "POSIXct" "POSIXt"
##
## $start_station_id
## [1] "integer"
##
## $start_station_name
## [1] "character"
##
## $start_station_latitude
## [1] "numeric"
##
## $start_station_longitude
## [1] "numeric"
##
## $end_station_id
## [1] "integer"
##
## $end_station_name
## [1] "character"
##
## $end_station_latitude
## [1] "numeric"
##
## $end_station_longitude
## [1] "numeric"
##
## $bikeid
## [1] "integer"
##
## $usertype
## [1] "character"
##
## $birth_year
## [1] "integer"
##
## $gender
## [1] "character"
```

Seems the column format is appropriate and no further change is needed at least in this stage.

Question: What is the trip duration distribution of Citibike trips?

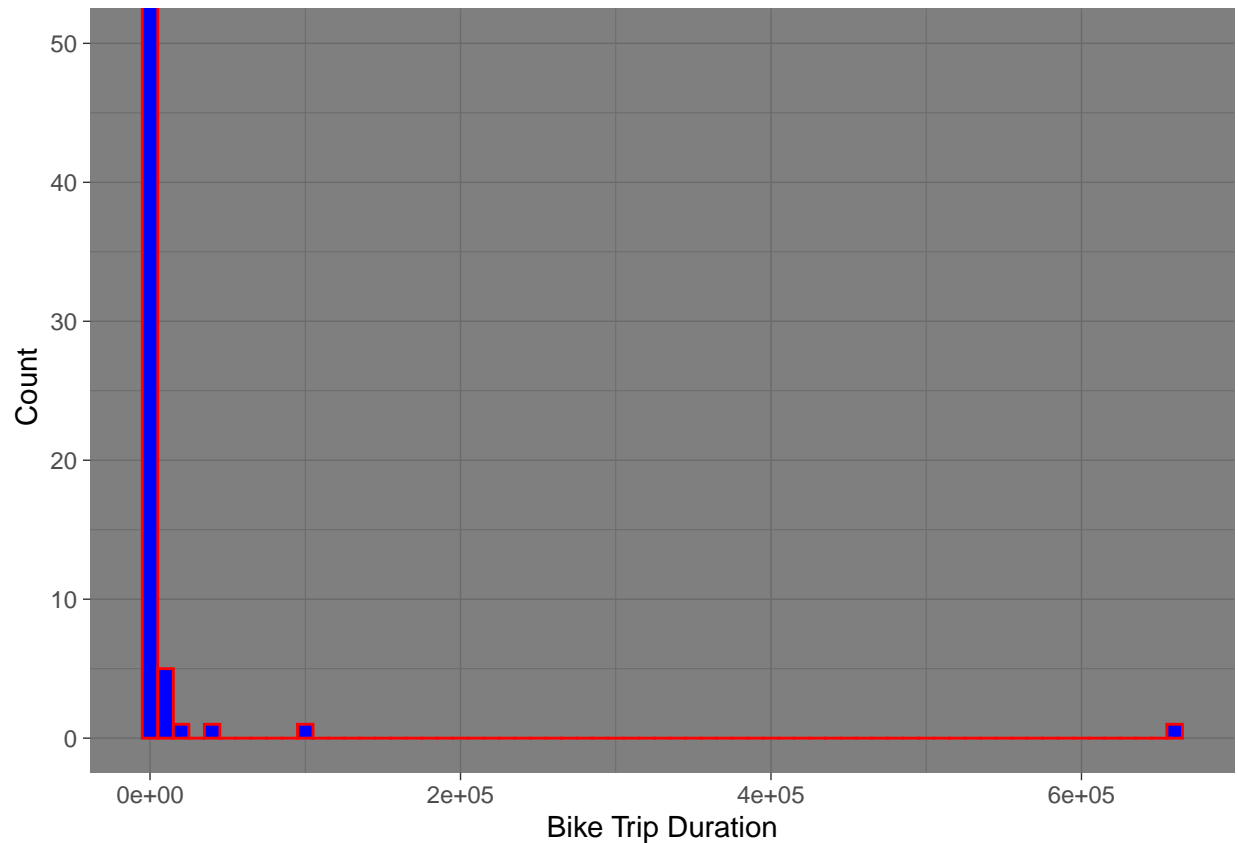
So, if we want to have look on the trip distribution, first we have look on the data-set we imported (1000) as the sample from the whole population. In the next steps, we further increase the size of the data to see is there any significant change from the analysis of the initial sample size.

```
ggplot(NY_sample, aes(tripduration)) +
  geom_histogram(fill = 'blue',color='red') +
  xlab("Bike Trip Duration") +
  ylab(" Count") +
  theme_dark()
```



Closer look on the *tripduration* we could figure out that the data has some outlier or possible some very off value. To have better understanding of this issue, we reduce the y axis limit to see where the outlier are:

```
ggplot(NY_sample, aes(tripduration)) +  
  geom_histogram(binwidth = 10000, fill = 'blue', color='red') +  
  coord_cartesian(ylim = c(0,50)) +  
  xlab("Bike Trip Duration") +  
  ylab(" Count") +  
  theme_dark()
```



Seems, there are some very long bike trip, let's have look on them; let's say the all bike trips more than 30000 sec:

```
NY_sample %>%
  filter(tripduration > 30000) %>%
  arrange(tripduration)
```

```
##   tripduration      starttime      stoptime start_station_id
## 1      36395 2016-03-15 21:36:55 2016-03-16 07:43:31        375
## 2      96235 2015-10-23 19:36:55 2015-10-24 22:20:51        493
## 3     659210 2015-11-04 07:20:10 2015-11-11 22:27:00        519
##   start_station_name start_station_latitude start_station_longitude
## 1 Mercer St & Bleecker St             40.72679             -73.99695
## 2      W 45 St & 6 Ave             40.75680             -73.98291
## 3  Pershing Square North             40.75187             -73.97771
##   end_station_id end_station_name end_station_latitude
## 1           3019 NYCBS Depot - DEL             40.71663
## 2           3019 NYCBS Depot - DEL             40.71663
## 3           3019 NYCBS Depot - DEL             40.71663
##   end_station_longitude bikeid  usertype birth_year gender
## 1             -73.98193  19910 Subscriber     1979   male
## 2             -73.98193  18137 Subscriber     1957   male
## 3             -73.98193  20894 Subscriber     1965   male
```

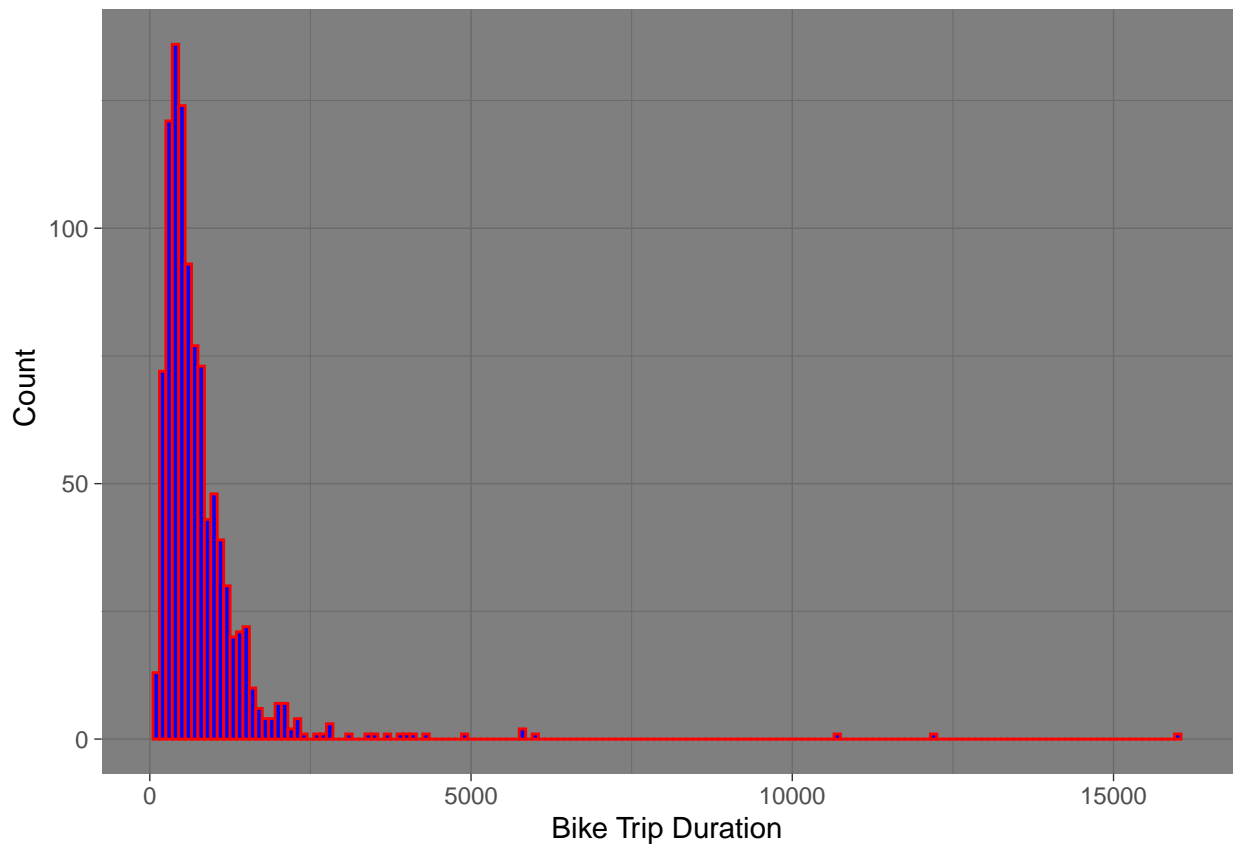
What we can say these three trips are about more than 10 hours of biking!, seems something unusual and probably are not the correct data due to error in registering the time or other issues, now we in this case

considering these cases, we do not include these three rows in the further analysis.

```
NY_sample_fil <- NY_sample %>%  
  filter(tripduration < 30000)
```

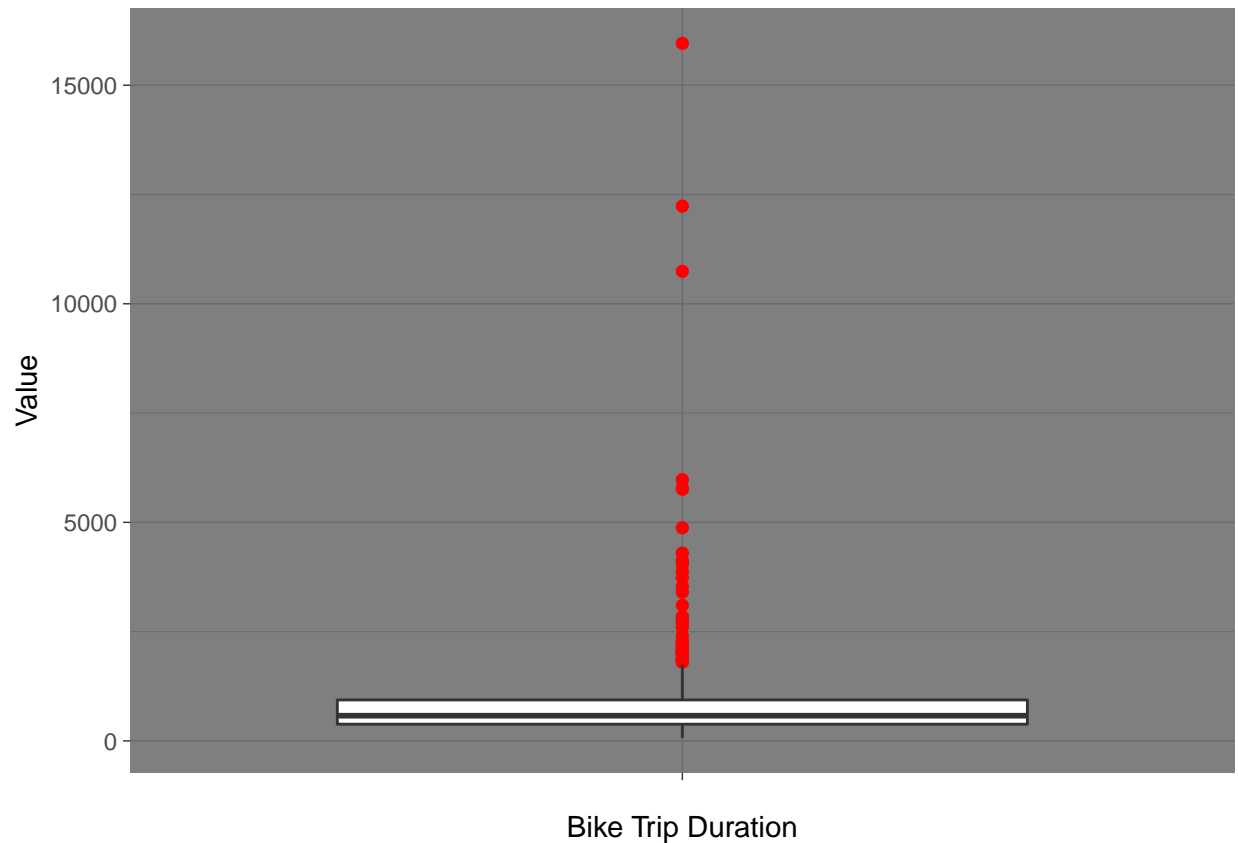
Now, let's plot the filtered data set and see the distribution of the *tripduration*:

```
ggplot(NY_sample_fil, aes(tripduration)) +  
  geom_histogram(binwidth = 100, fill='blue', color='red') +  
  xlab("Bike Trip Duration") +  
  ylab("Count") +  
  theme_dark()
```



Ok, now we could see that the data has the outlier. To further check this issue, we draw the BOX plot of the data to look:

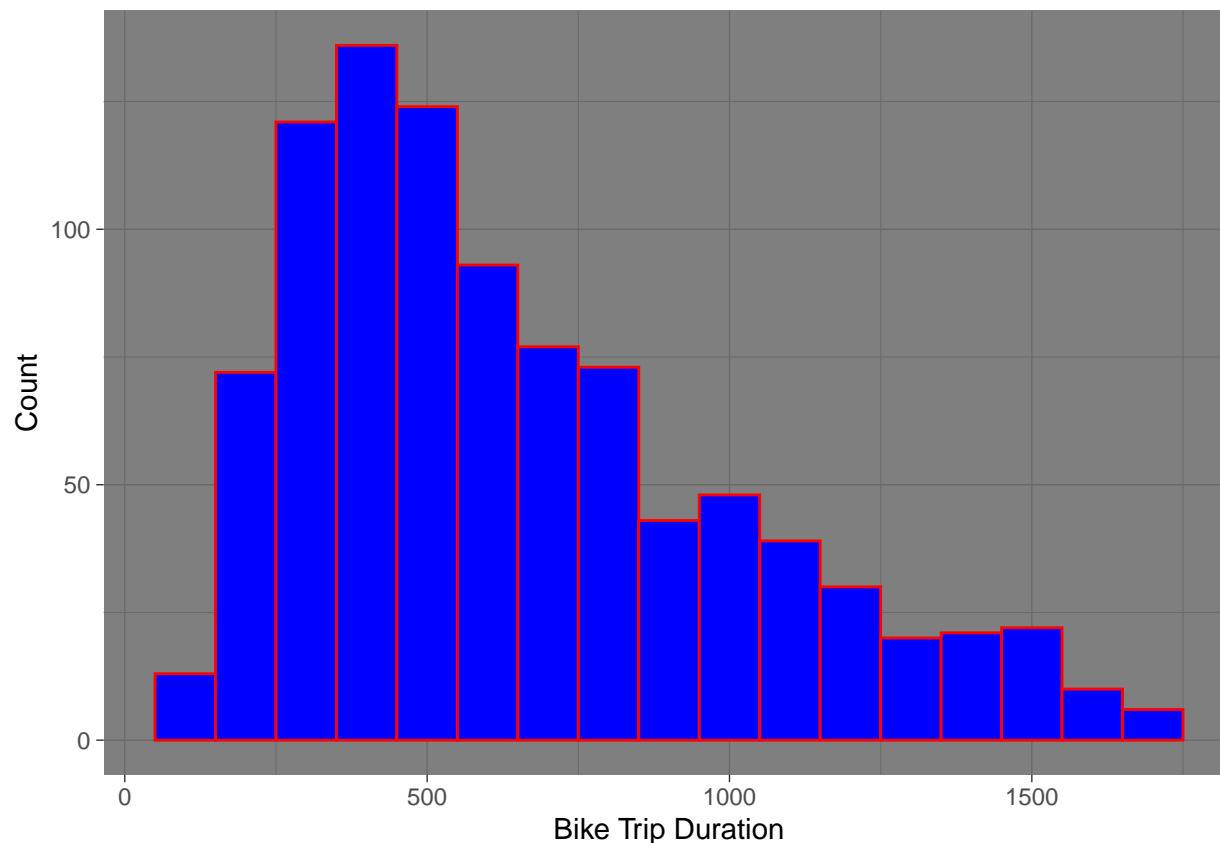
```
ggplot(NY_sample_fil, aes(x = "", y=tripduration)) +  
  geom_boxplot(outlier.colour="red",  
               outlier.shape=16,  
               outlier.size=2, notch=FALSE) +  
  xlab("Bike Trip Duration") +  
  ylab("Value") +  
  theme_dark()
```



Now, to be more clear about the real distribution of trip time, we again plot the distribution, yet this time we do not include the outlier in the dataset. By definition, here the outlier means the data *outside* of the 1.5 IQR from the 75% and 25 % quantile of the data. Have look on the distribution without the outliers:

```
quantile_query <- quantile(NY_sample_fil$tripduration,c(0.25,0.75))
lower_range <- quantile_query[1]-1.5*IQR(NY_sample_fil$tripduration)
upper_range <- quantile_query[2]+1.5*IQR(NY_sample_fil$tripduration)

NY_sample_fil %>%
  filter(tripduration > lower_range & tripduration < upper_range) %>%
  ggplot(aes(tripduration)) +
  geom_histogram(binwidth = 100, fill='blue', color='red') +
  xlab("Bike Trip Duration") +
  ylab("Count") +
  theme_dark()
```

However, the above data analysis were for the small sample (10000 rows), how about the large data-set. Here, we include the 1,000,000 rows to be included in the analysis to provide the better understanding the distribution.

```
# Project ID
project_id <- "xenon-muse-261616" # put your project ID here

# query
sql_string_tot <- "SELECT
  tripduration
FROM
  [bigquery-public-data.new_york.citibike_trips]
LIMIT
  1000000"
# Execute the query and store the result
NY_Lsample <- query_exec(sql_string_tot, project = project_id,
                        useLegacySql = FALSE, max_pages = 100)
```

```
## 0 bytes processed
```

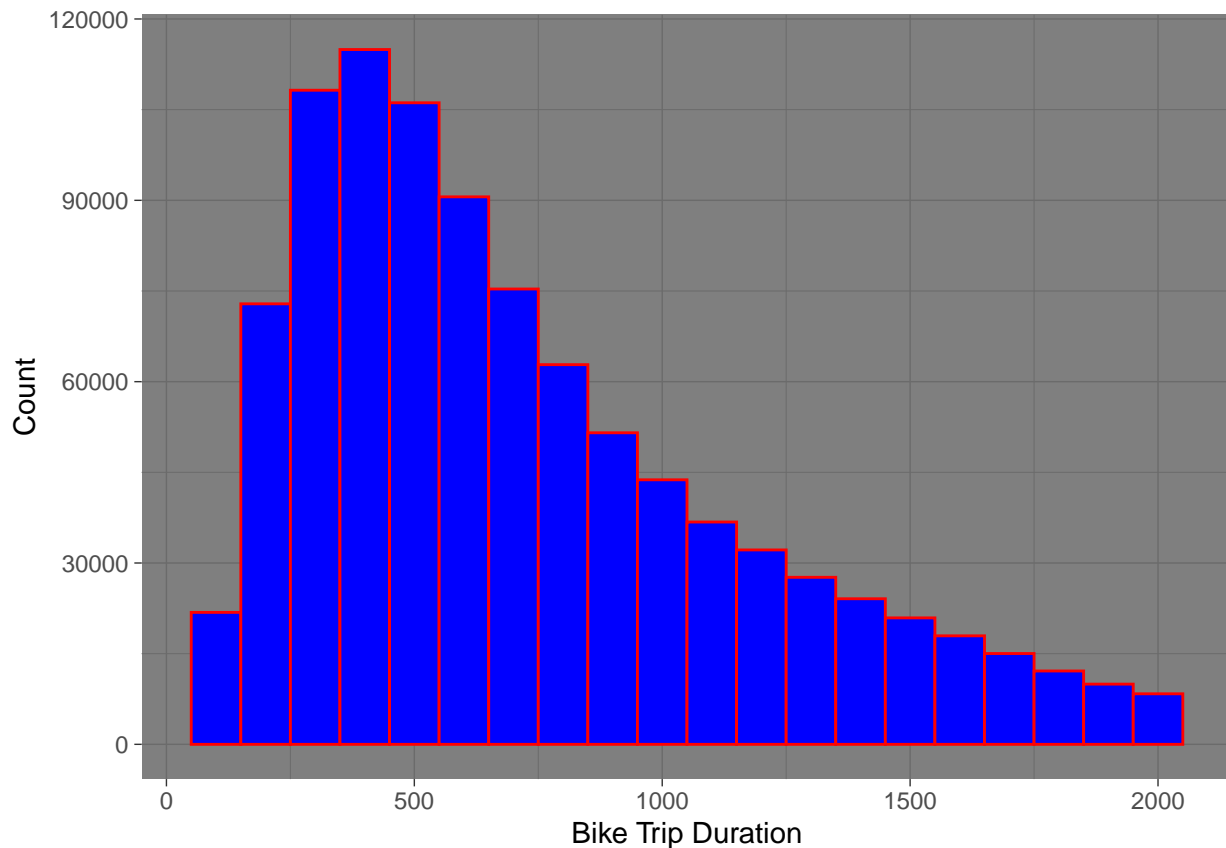
We again plot the data-set this time without the outlier when the larger rows were included in the data-set.

```
quantile_queryL <- quantile(NY_Lsample$tripduration,c(0.25,0.75))
lower_rangeL <- quantile_queryL[1]-1.5*IQR(NY_Lsample$tripduration)
upper_rangeL <- quantile_queryL[2]+1.5*IQR(NY_Lsample$tripduration)
```

```

NY_Lsample %>%
  filter(tripduration > lower_rangeL & tripduration < upper_rangeL) %>%
  ggplot(aes(tripduration)) +
  geom_histogram(binwidth = 100, fill='blue', color='red') +
  xlab("Bike Trip Duration") +
  ylab("Count") +
  theme_dark()

```



Question: What is the most popular Citibike trip?

So, from the previous question we got to know the data and how it looks like. For this question we only import the two rows we needed in the question, *start_station_id* and *end_station_id*, also we get the about 1,000,000 rows this time:

```

# Project ID
project_id <- "xenon-muse-261616" # put your project ID here

# query
sql_string_tot_trip <- "SELECT
  start_station_id,
  end_station_id
FROM
  [bigquery-public-data.new_york.citibike_trips]"

```

```

LIMIT
1000000"
# Execute the query and store the result
NYL_trip <- query_exec(sql_string_tot_trip, project = project_id, useLegacySql = FALSE, max_pages = 100)

```

```
## 0 bytes processed
```

have the brief look on the data:

```
head(NYL_trip)
```

```

##   start_station_id end_station_id
## 1                512            512
## 2                512            512
## 3                512            512
## 4                512            512
## 5                512            512
## 6                512            512

```

Here, we group the data based on the *start_station_id* and as well *end_station_id* and then we count the number of trip between any two starting and end station:

```

NYL_trip_group <- NYL_trip %>%
  count(start_station_id, end_station_id) %>%
  filter(n > 300) %>%
  arrange(desc(n))
NYL_trip_group

```

```

## # A tibble: 10 x 3
##   start_station_id end_station_id     n
##   <int>           <int> <int>
## 1       2006         2006  1615
## 2        281         281   646
## 3        363         327   508
## 4        514         426   437
## 5        426         426   383
## 6        514         514   359
## 7        435         462   354
## 8        499         499   335
## 9        116         521   332
## 10       457         457   325


```

Now, we can plot the result in one river plot which thickness of the each river show the *count* of the trip between two two stations:

```

library(alluvial)
alluvial(NYL_trip_group[,1:2], freq=NYL_trip_group$n, xw=0.0, alpha=0.8,
  gap.width=0.1, col= 1:10, border="white")


```



However, maybe the trip between the same *station_id* may not be informative, therefore in the new plot we removed the same start and end stations and includes the only trips where have different start and end point:

```
library(alluvial)
NYL_trip_group_nonsim <- NYL_trip %>%
  filter(start_station_id != end_station_id) %>%
  count(start_station_id, end_station_id) %>%
  filter(n > 300) %>%
  arrange(desc(n))

alluvial(NYL_trip_group_nonsim[,1:2], freq=NYL_trip_group_nonsim$n, xw=0.0, alpha=0.8,
  gap.width=0.1, col= 1:4, border="white")
```



Question: Were there new bike stations introduced or removed at any point in time? What makes you think it were or weren't?

So, to answer this question we first need to import the two columns , *starttime* and *start_station_id*. Here we assume if new station was introduced in the start station so that the new trips could be done using this station.

```
# Project ID
project_id <- "xenon-muse-261616" # put your project ID here

# query
sql_string_newsta <- "SELECT
  starttime,
  start_station_id
FROM
  [bigquery-public-data.new_york.citibike_trips]
LIMIT
  1000000"
# Execute the query and store the result
NYL_newsta <- query_exec(sql_string_newsta, project = project_id,
  useLegacySql = FALSE, max_pages = 100)
```

```
## 0 bytes processed
```

Now, the code below shows the distinct start station id per month. So, this could be helpful when in the next stage when we want to know which station was added or removed.

```
NYL_newsta_group <- NYL_newsta %>%  
  mutate(month = as.numeric(format(starttime, "%m")), year = as.numeric(format(starttime, "%Y")),  
         day=as.numeric(format(starttime, "%d"))) %>%  
  select(year, month, start_station_id) %>%  
  distinct(year, month, start_station_id, .keep_all = TRUE) %>%  
  group_by(year, month, start_station_id) %>%  
  arrange(year, month)
```

For example, we could ask what station was added to the trips on month 8 of 2013. To do that, we compare the whole recorded trip until the month 7 of 2013 and compare it with the month 8 of 2013.

```
old <- NYL_newsta_group %>%  
  filter(year<=2013 & month <= 10) %>%  
  pull(start_station_id)  
  
old <- unique(old)  
  
new <- NYL_newsta_group %>%  
  filter(year==2014 & month == 10) %>%  
  pull(start_station_id)
```

The new stations added to the biking trips on the month 8 of 2013 are includes:

```
new[!(new %in% old)]
```

```
## [1] 519 530 320 491
```

Now, we can see that the station number 502 was added to the traveling route on month 8th of 2013.

Now, let's look what station was removed during the time. Let's say what station were removed at October of 2014 when compared to October of the 2013.

```
old <- NYL_newsta_group %>%  
  filter(year<=2013 & month <= 10) %>%  
  pull(start_station_id)  
  
old <- unique(old)  
  
new <- NYL_newsta_group %>%  
  filter(year==2014 & month == 10) %>%  
  pull(start_station_id)  
  
old[!(old %in% new)]
```

```
## [1] 517 243
```

The stations named 517 and 243 were available on October of 2013 but, were not available on the October of 2014.

For the number of stations added to the route, we here calculate the number of stations were added at 2015, per month.

```

new_station <- numeric(12)
for (mnt in 1:12) {
  if (mnt==1) {
    yr = 2014
    mnth = 12
    mnt = 1
  } else {
    yr = 2015
    mnth = mnt-1
  }

  old <- NYL_newsta_group %>%
    filter(year<=yr & month <= mnth) %>%
    pull(start_station_id)

  old <- unique(old)

  new <- NYL_newsta_group %>%
    filter(year==2015 & month == (mnt)) %>%
    pull(start_station_id)

  new_station[mnt] <- length(new[!(new %in% old)])
}

months <- c("January", "February", "March", "April", "May", "June", "July",
            "August", "September", "October", "November", "December" )

df_newstat <- data.frame(month=months,
                        New_Stations = new_station)

ggplot(data=df_newstat, aes(x=factor(month,levels=month.name), y=New_Stations)) +
  geom_bar(stat="identity", fill='blue', color='red') +
  xlab("Month, at 2015") +
  ylab("Number of new stations") +
  theme_dark()

```

